

MEMORY MANAGEMENT IN

OS → *Long Session*

SESSION 6

LINK IN THE DESCRIPTION FOR ALL SESSIONS

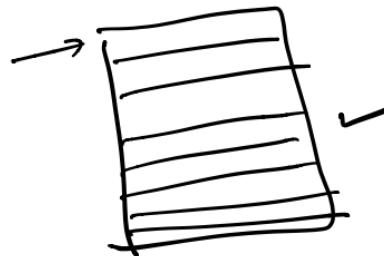
PREVIOUS SESSION

- Process synchronization ✓
- Race condition ✓
- Critical section problem ✓
- Software and hardware based solutions ✓
- Semaphore and mutex ✓
- Producer consumer problem ✓
- Reader writer problem ✓

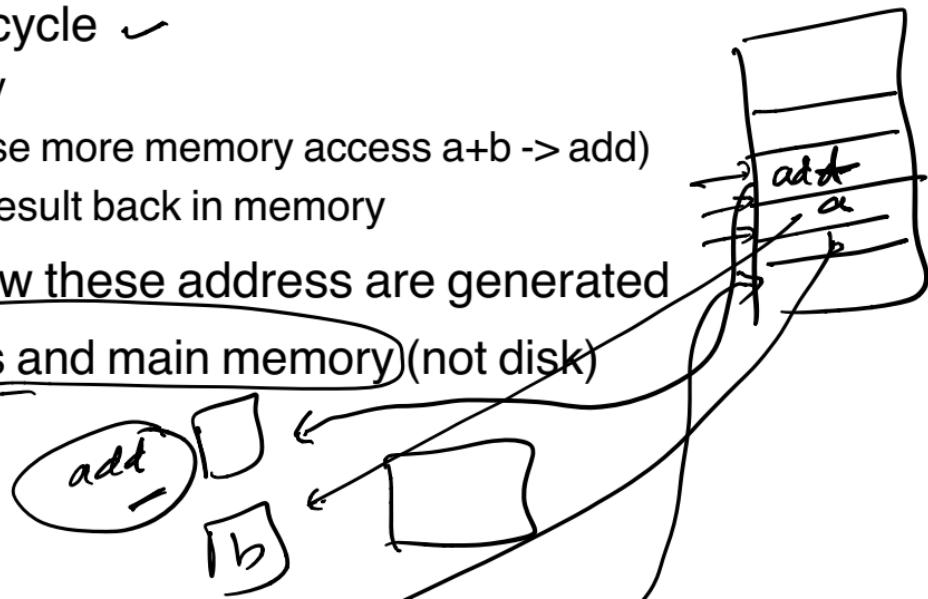
INTERVIEW QUESTIONS

- Difference between 32-bit and 64-bit system ✓
- Logical memory vs Physical memory ✓
- Normal paging vs demand paging vs pure demand paging ✓
- Segmentation and paging ✓
- Virtual memory and its advantage ✓
- Paging vs segmentation ✓
- Physical address vs virtual address ✓
- Inverted paging ✓

MAIN MEMORY



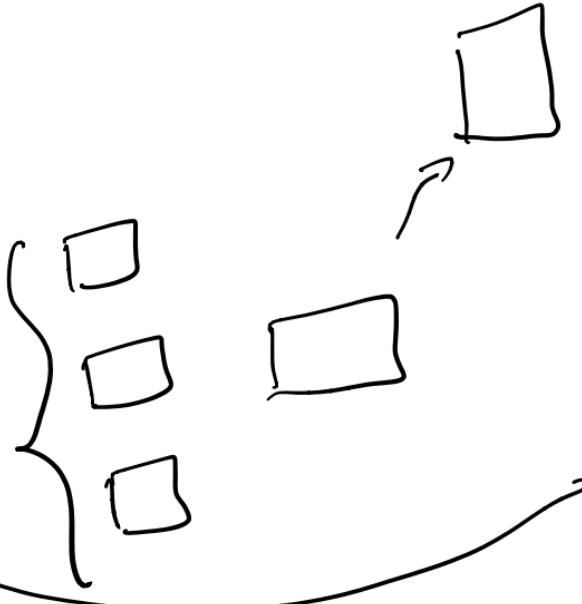
- Memory consists of a large array of words or bytes, each with its own address.
- Typical instruction-execution cycle ✓
 - Fetch instruction from memory
 - Instruction decoding (may cause more memory access $a+b \rightarrow \text{add}$)
 - Execute instruction and store result back in memory
- Memory unit doesn't know how these address are generated
- CPU can only access registers and main memory (not disk)



FASTEST MEMORY

- Registers ✓
- Processor Cache – L1, L2 and L3 ✓
- Main Memory – RAM
- Disk

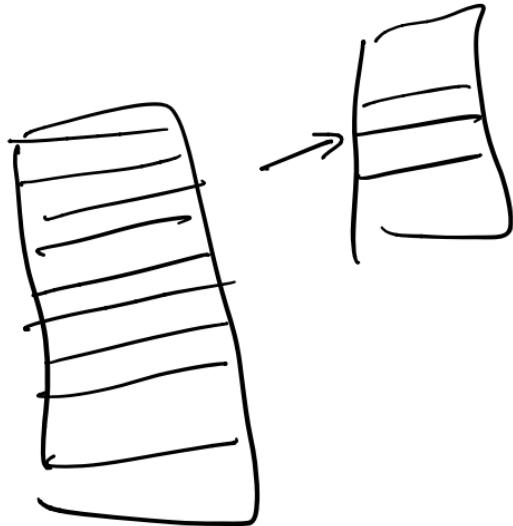
c
atb



ADDRESS BINDING



- Compile time ✓
 - Starting address is known ✓
 - Starting location changes ? ✓

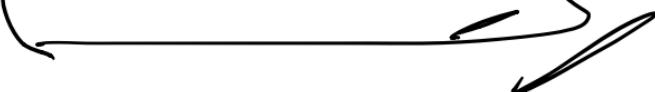


- Load time
 - Address not known at compile time ==
 - Final binding is delayed till load time ✓
 - Relocatable code



- Execution time ✓
 - If process can be moved during execution? ✓
 - Binding is delayed till run time ✓

Used in most modern OS



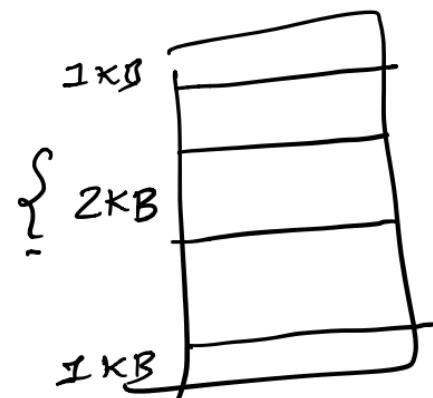
gcc



CONTIGUOUS MEMORY ALLOCATION

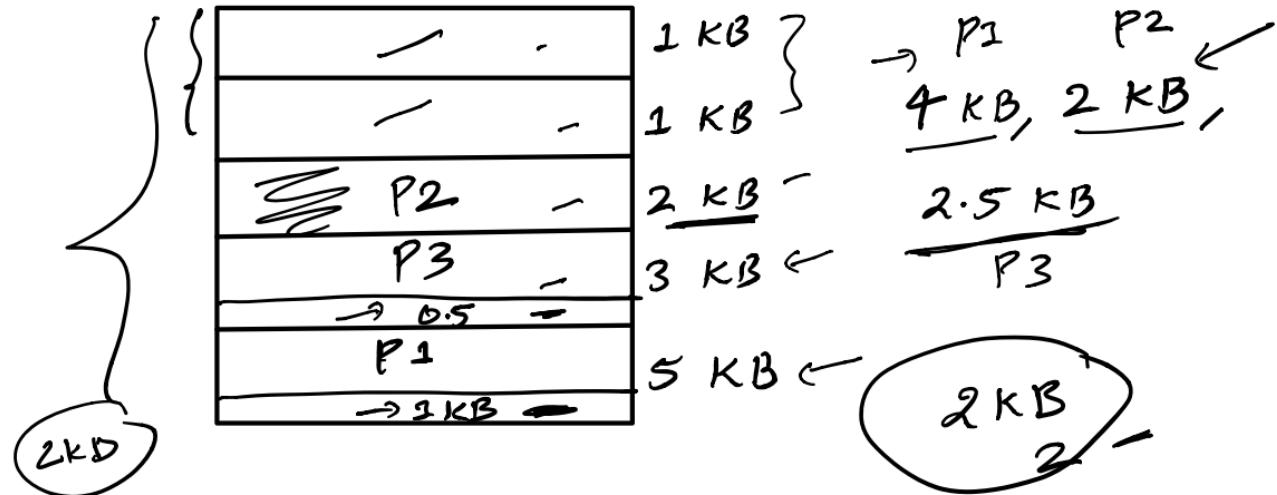
- Traditional approach ✓
- Allocate contiguous memory blocks for a process
- Two approaches ✓
 - Fixed partitioning ✓
 - Variable partitioning ✓

P_1 2 KB



FIXED PARTITIONING

- Memory is divided into partitions of fixed size which cannot change
- Each partition can only contain one process
- Find a suitable partition and allocate it



DISADVANTAGE

- Degree of multiprogramming bounded by number of partitions
- Size of process is bounded by size of largest partition
- Suffers from internal fragmentation

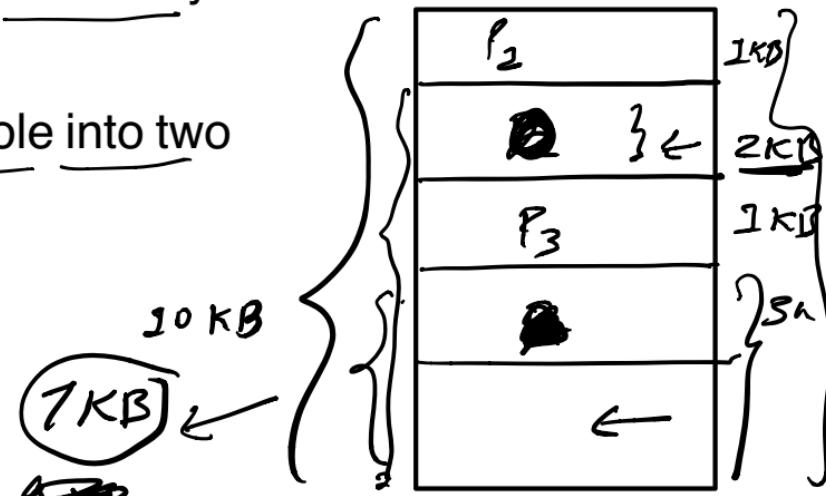


VARIABLE PARTITION METHOD

- Partitions are created dynamically as process are loaded
- Free partitions -> holes
- Initially, one hole whose size is of main memory
- OS keeps tracks of these holes
- If process size > selected hole, split hole into two
 - Give one to process
 - Other is returned to the OS
- Process terminates, hole is released.
- Two adjacent holes are merged

2 KB, 6 KB

$P_1 \rightarrow 1 \text{ KB}$
 $P_2 \rightarrow 2 \text{ KB}$
 $P_3 \rightarrow 1 \text{ KB}$
 $P_4 \rightarrow 3 \text{ KB}$

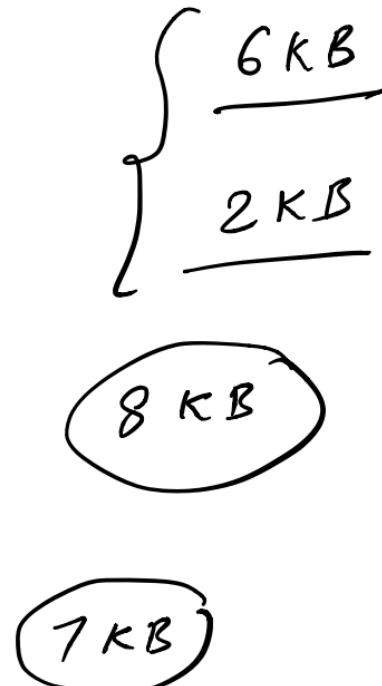


~~7 KB~~ ~~6 KB~~ ~~2 KB~~, 2 KB, ~~2 KB~~

- No internal fragmentation ✓
- Degree of multiprogramming ✓

Disadvantage ✓

- Size of the process is limited by size of biggest hole
- Leads to external fragmentation



ALGORITHMS FOR VARIABLE PARTITIONING

- FIRST FIT ✓

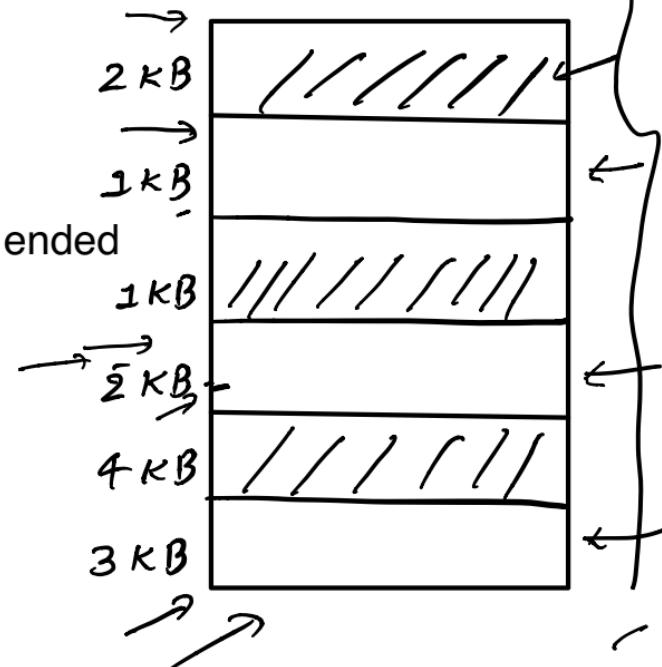
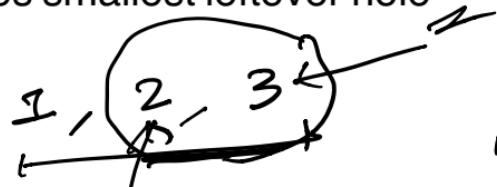
- Scan from beginning
- Allocate the first hole that is big enough

- Next Fit ✓

- Scan from location at which previous searched ended
- Allocate the first hole that is big enough

- BEST FIT ✓

- Allocate the smallest hole that is big enough
- Produces smallest leftover hole



Selected \geq process Size

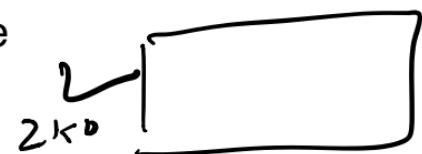
$\xrightarrow{P} \underline{2KB}$

- WORST FIT ✓

- Allocate the largest hole that is big enough
- Produces largest leftover hole
- Leftover hole may be more useful than best-fit created hole



$2KB$



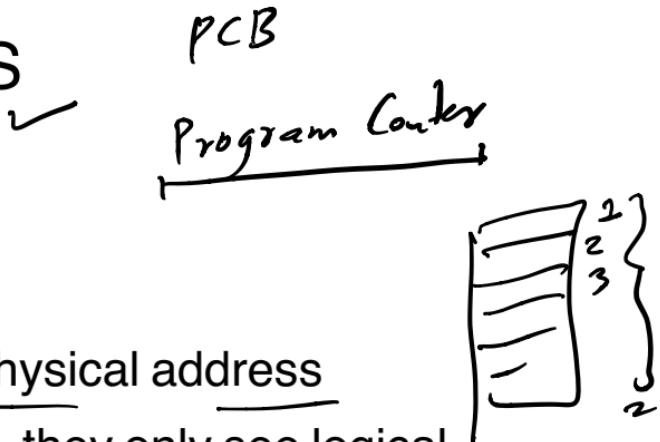
$2KB$

Example - <https://youtu.be/fz9v-SZt3cQ>

$\xrightarrow{2}$

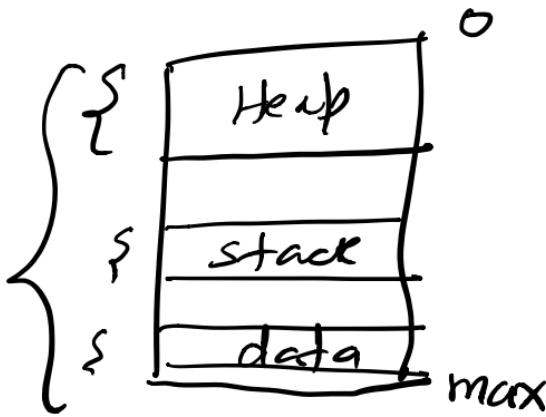
LOGICAL VS PHYSICAL ADDRESS

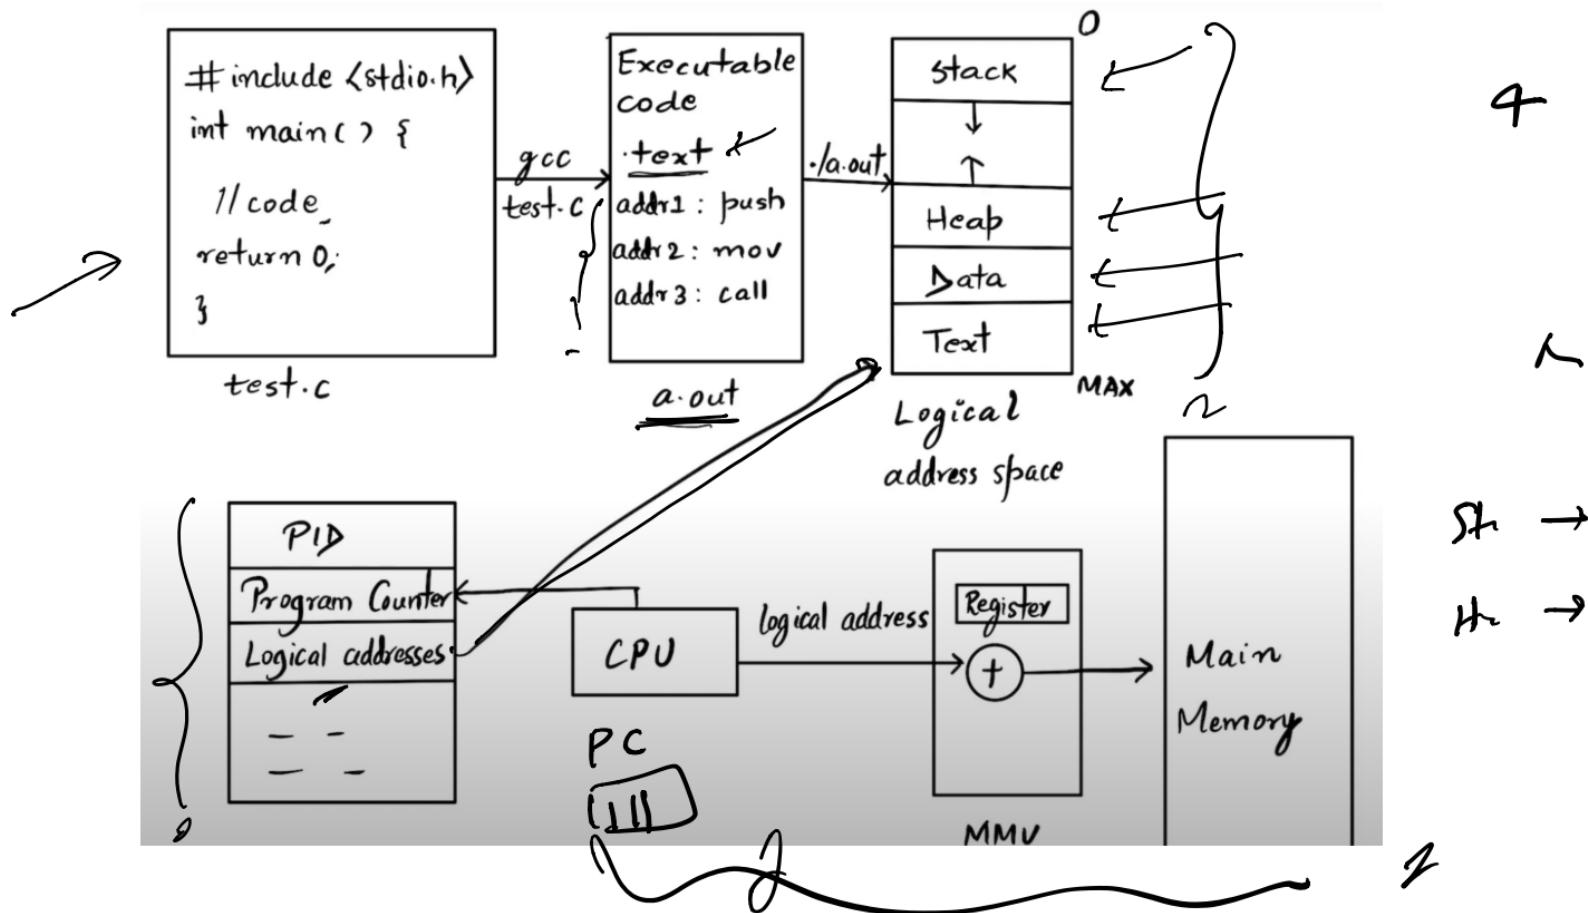
- Address which the CPU sees -> logical address
- Address contained in the program counter
- Address that exists in main memory is called physical address
- User programs never see the physical address, they only see logical address



addr = malloc(int)

Print("%p", &addr) -> logical or physical ?

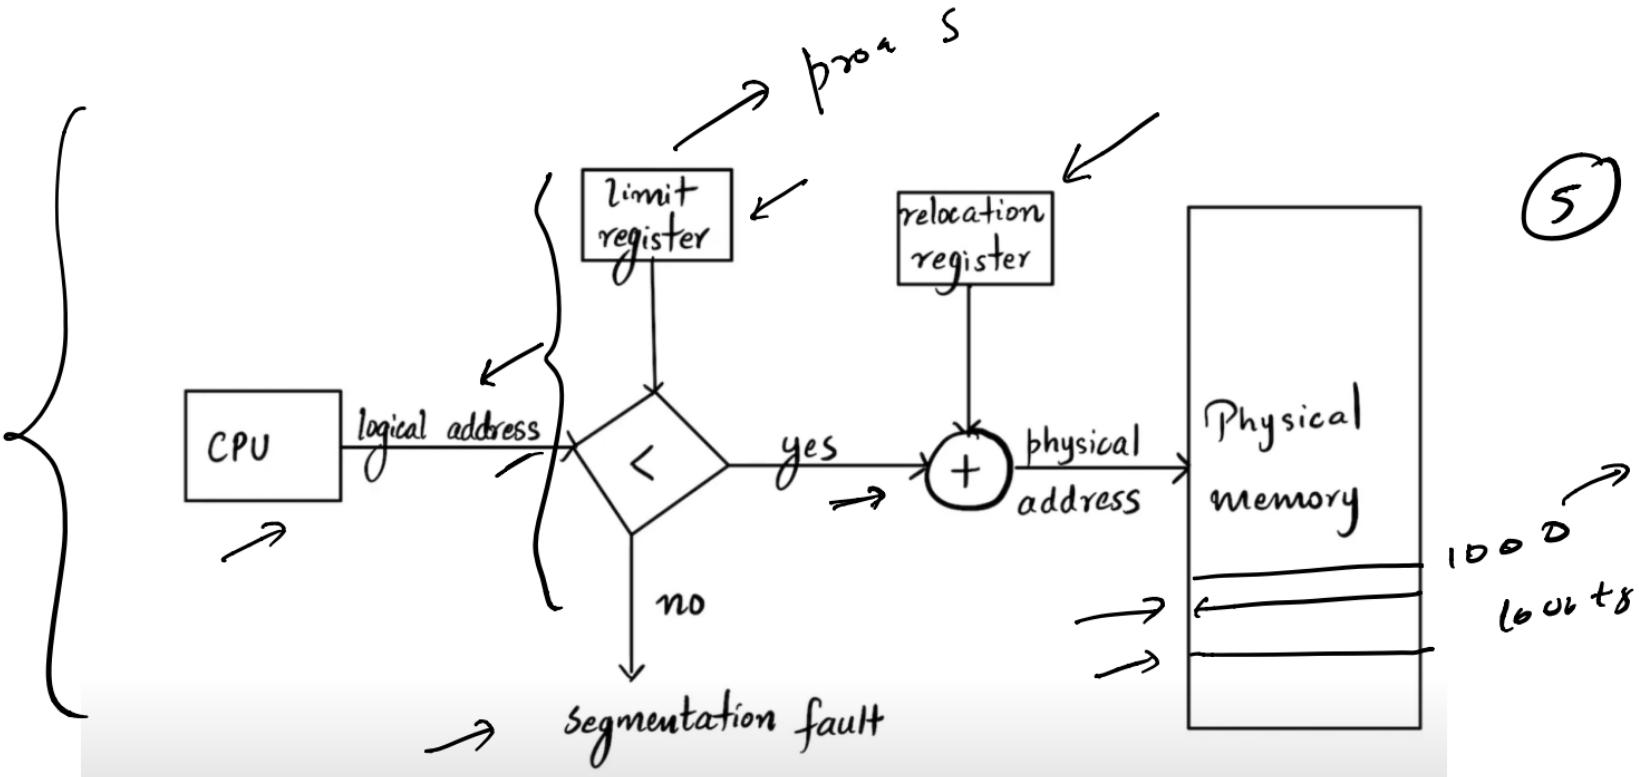




S → -/-
H → --
T → --

MEMORY MAPPING AND PROTECTION IN CONTIGUOUS MEMORY ALLOCATION

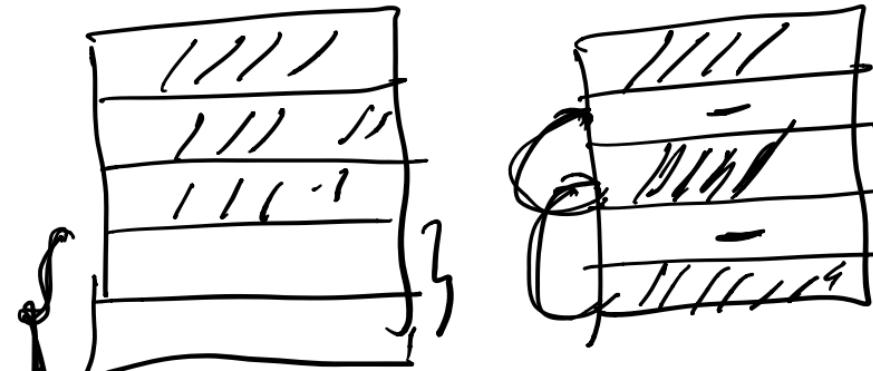
- The logical address need to be converted into physical address for instruction execution
- Also, one process should not be able to access other process memory
- Use limit register and relocation register ✓
- Limit register ≠ size of the process
- Relocation register = Starting physical address of the process



NONCONTIGUOUS MEMORY ALLOCATION

TECHNIQUES

- Limitation with contiguous memory allocation ✓
 - Internal fragmentation – Fixed partition ↗
 - External fragmentation – Variable partition ↗
- Use compaction to solve the external fragmentation ↗
- But compaction is very expensive because of copying operation (I/O)
- Other solutions?



- Allow process to have noncontiguous physical address space
- Break down logical address of a process into chunks and allocate them
- Noncontiguous memory management techniques
 - Paging ✓
 - Segmentation ✓
 - Paging + Segmentation ✓

PAGING TECHNIQUE

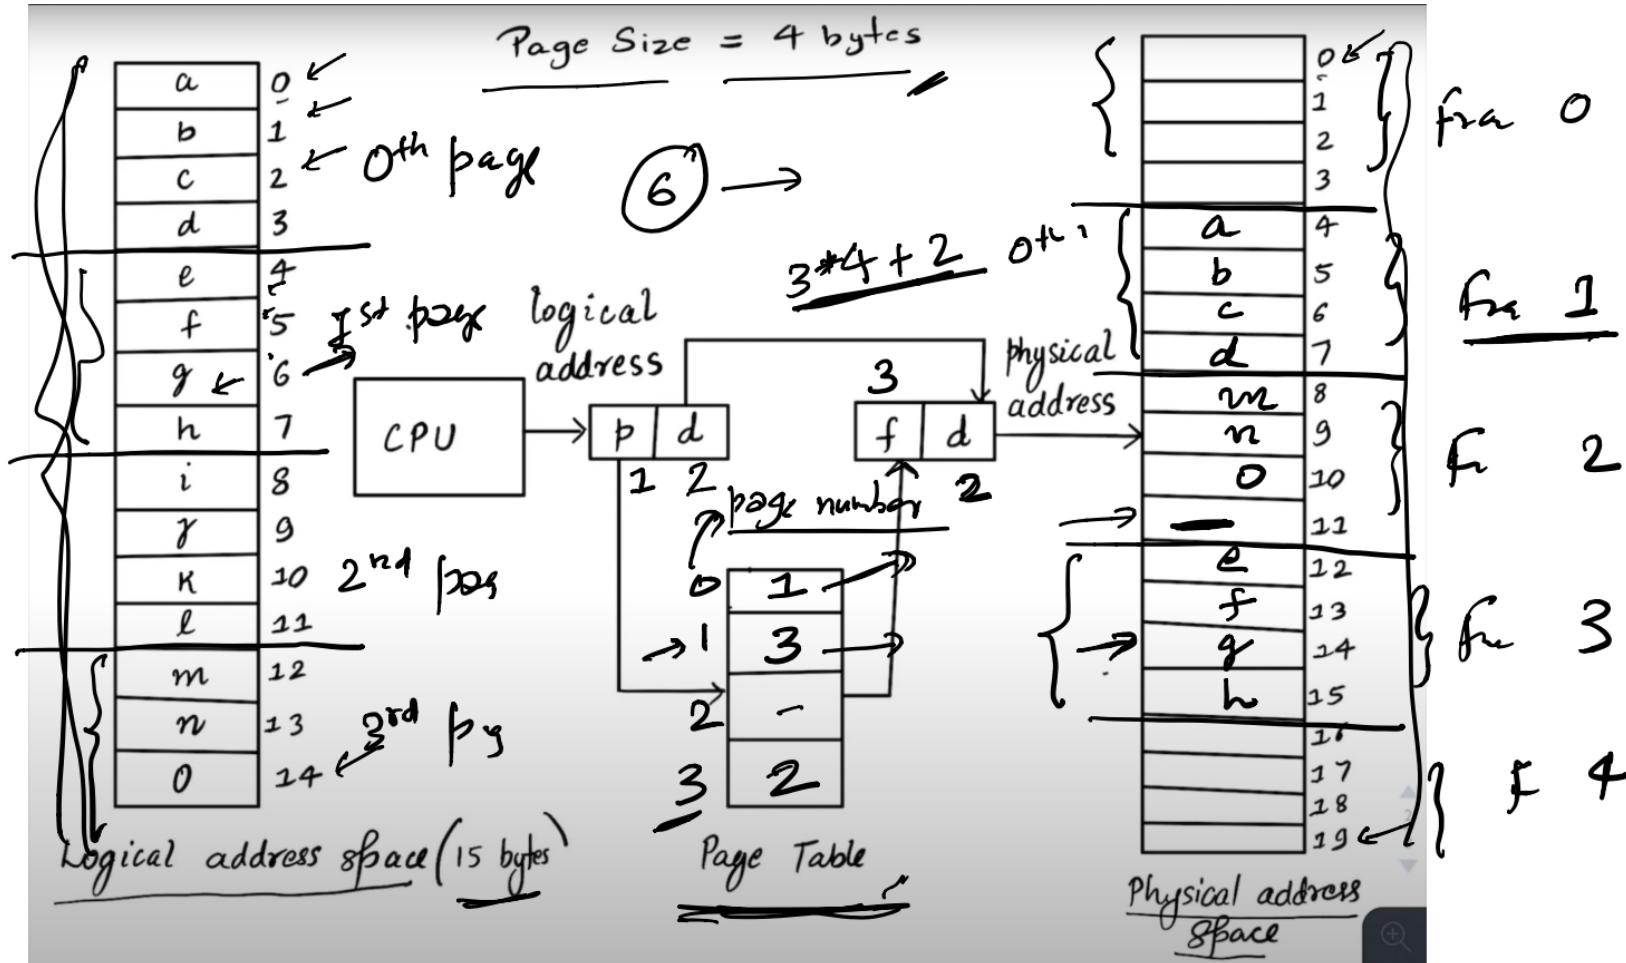
- Assume byte addressable memory
- Each byte is a unique address



- Break down the logical memory into partitions called pages
- Break down the physical memory into partitions called frames
- Page size = frame size
- Use page tables to map logical address to physical address
- Note: In simple paging, all the pages are brought in memory before process starts executing.

N pages

- If there are N pages, then at least N frames must be available for process to run. See demand paging
- Logical address – page number + page offset
- MMU combines page table mapping with page offset to arrive at physical address



$$\underline{3 * 4 + 2}$$

- Page number = logical address / page size
- Page offset = logical address % page size
- Can be simplified if page size is power of 2
- Number of pages = logical memory size / page size
- Number of frames = physical memory size / frame size
- Frame size = page size
- Avoids external fragmentation
- Some internal fragmentation (Fixed size partitioning)
- Same as using a table of base/relocation registers

$$\frac{6}{4} = \underline{\underline{1}}$$

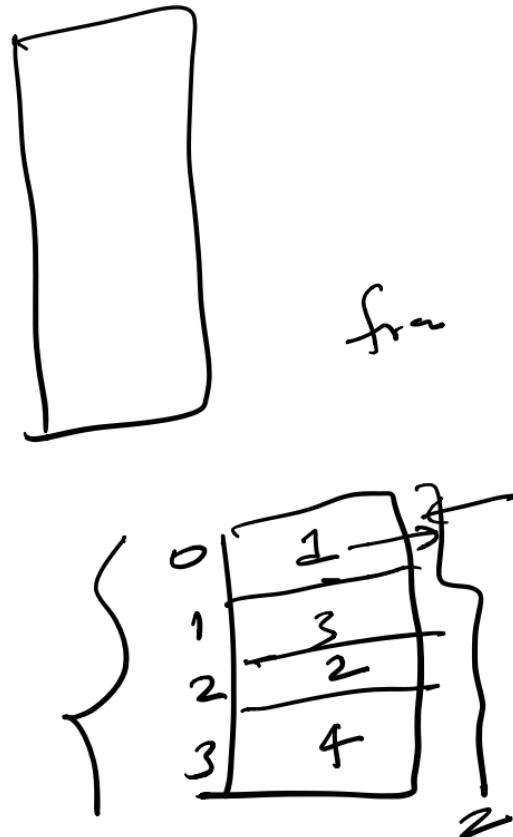
$$\frac{15}{4} = \underline{\underline{3..}}$$

$$\frac{20}{4} = \underline{\underline{5}}$$

CLEARLY UNDERSTAND ✓

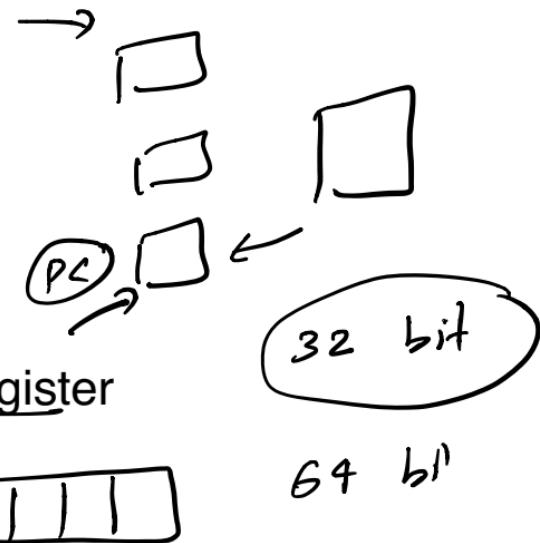
- PAGE ✓
- FRAME ✓
- PAGE NUMBER ✓
- FRAME NUMBER ✓
- PAGE SIZE ✓
- FRAME SIZE ✓
- LOGICAL ADDRESS SPACE ✓
- PHYSICAL ADDRESS SPACE
- PAGE TABLE ✓
- PAGE TABLE SIZE ✓

4* -



32-BIT VS 64-BIT SYSTEMS

- 32-bit -> CPU has 32 bit registers
- 64-bit -> CPU has 64 bit registers
- Logical address is stored in program counter register
- 32 bit -> max logical address size -> 32 bit
- Maximum address generated with 32 bit = 2^{32} = 4GB (binary basics)
- You cannot run a 5GB game on 32-bit processor system



34

32 bit → 2^{32}

PAGE SIZE AS POWER OF 2

- Page size = 2^n ✓
- Logical address space = 2^m (m-bit system) ✓

Binary basics: ✓

- N bits → 2^N different values from 0 to $2^{(N-1)}$
- 2^N values → N bits

$$\begin{array}{ccccccc} & & & & 3 \text{ bi} & & \\ 8 & \xrightarrow{\quad} & 3 & & & & \\ & & \log_2 8 \rightarrow 3 & & 0-1 & \rightarrow & 8 \\ & & & & & & \end{array}$$

- Total pages = $2^m / 2^n = \underline{2^{(m-n)}}$ pages = 0 to $2^{(m-n)-1}$
- How many bits you need to represent $2^{(m-n)}$ page numbers?
- (m-n) bits to represent page number
- Logical address = page number + page offset
- $m = (m-n) + \text{page offset} \rightarrow \text{page offset size} = \underline{n \text{ bits}}$
- Page offset can vary from 0 to 2^N-1
- Similar approach can be followed for frames

$$2^{\chi}$$



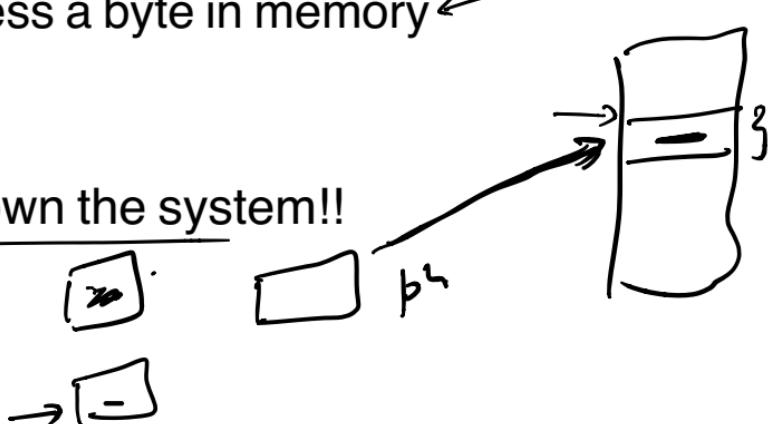
$$(2^{(m-n)} - 1)$$

logical address / pg

$$2^m / 2^n$$

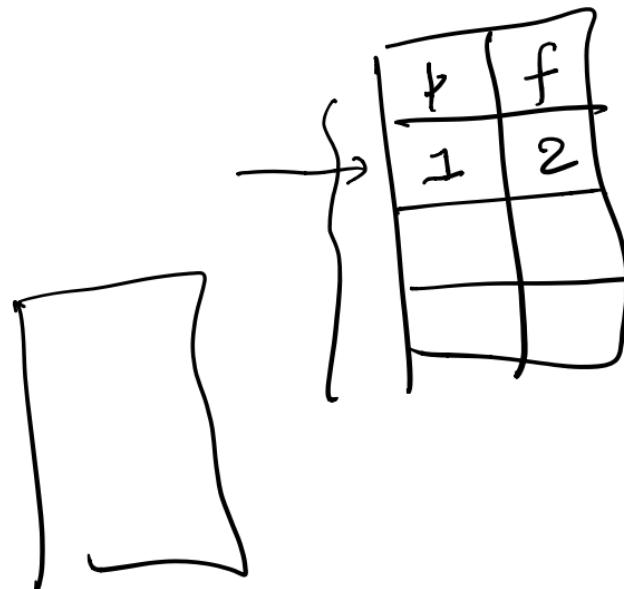
ISSUE WITH KEEPING PAGE TABLE IN MAIN MEMORY

- The page table is kept in main memory, and a points to the page table. Changing page tables requires changing only this one register, substantially reducing context-switch time
- Threads vs Process context switch time (Thread lecture)
- Two memory access required to access a byte in memory
 - Page table access
 - Access the required byte
- Doing this for every byte will slow down the system!!



TRANSLATION LOOK-ASIDE BUFFER (TLB)

- A small, expensive cache that sits close to CPU
- Key, value data structure ✓
- Key -> page number ✓
- Value -> Frame number ✓
- Total entries = 64-1024 ✓
- TLB hit and TLB miss



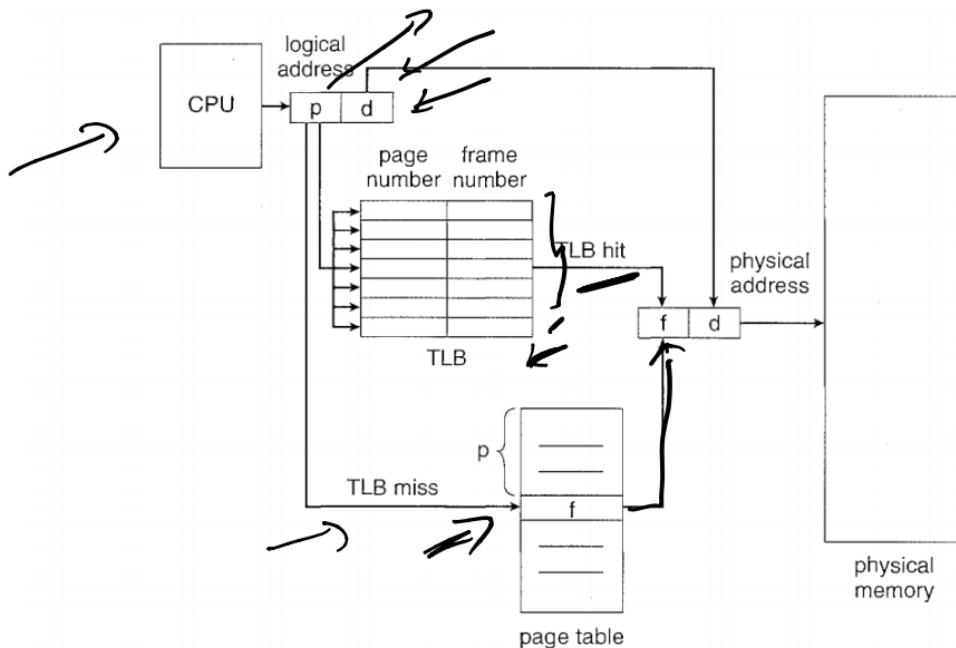


Figure 8.11 Paging hardware with TLB.

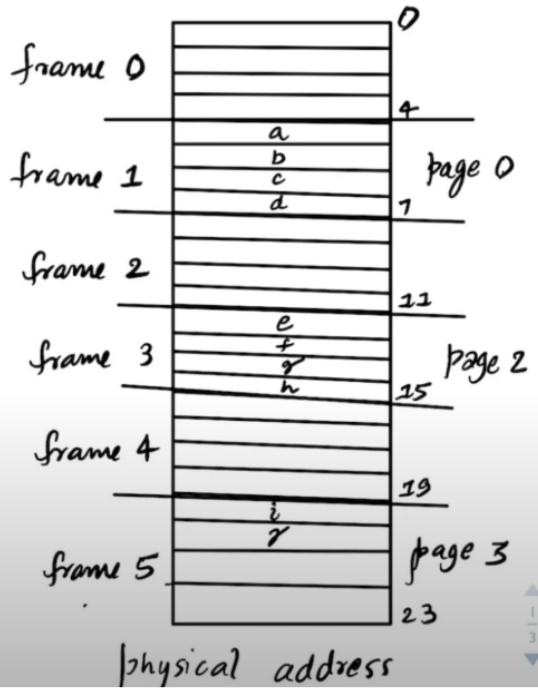
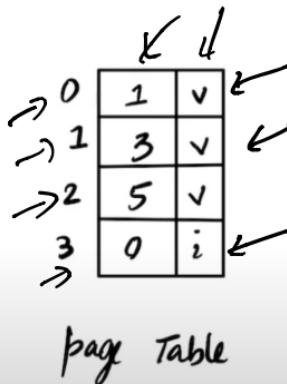
MEMORY PROTECTION IN PAGING

- Protection bits can be associated with along side frame number in page table
- Valid-invalid bit -> page exists in logical memory or not
- Dirty bit -> Page has been modified or not
- Cache enabled/disabled -> where caching is enabled for page or not
- Reference bit -> whether page was referenced in last clock cycle

Assume that the System has 4-bit logical address space.

page size = 4 bit

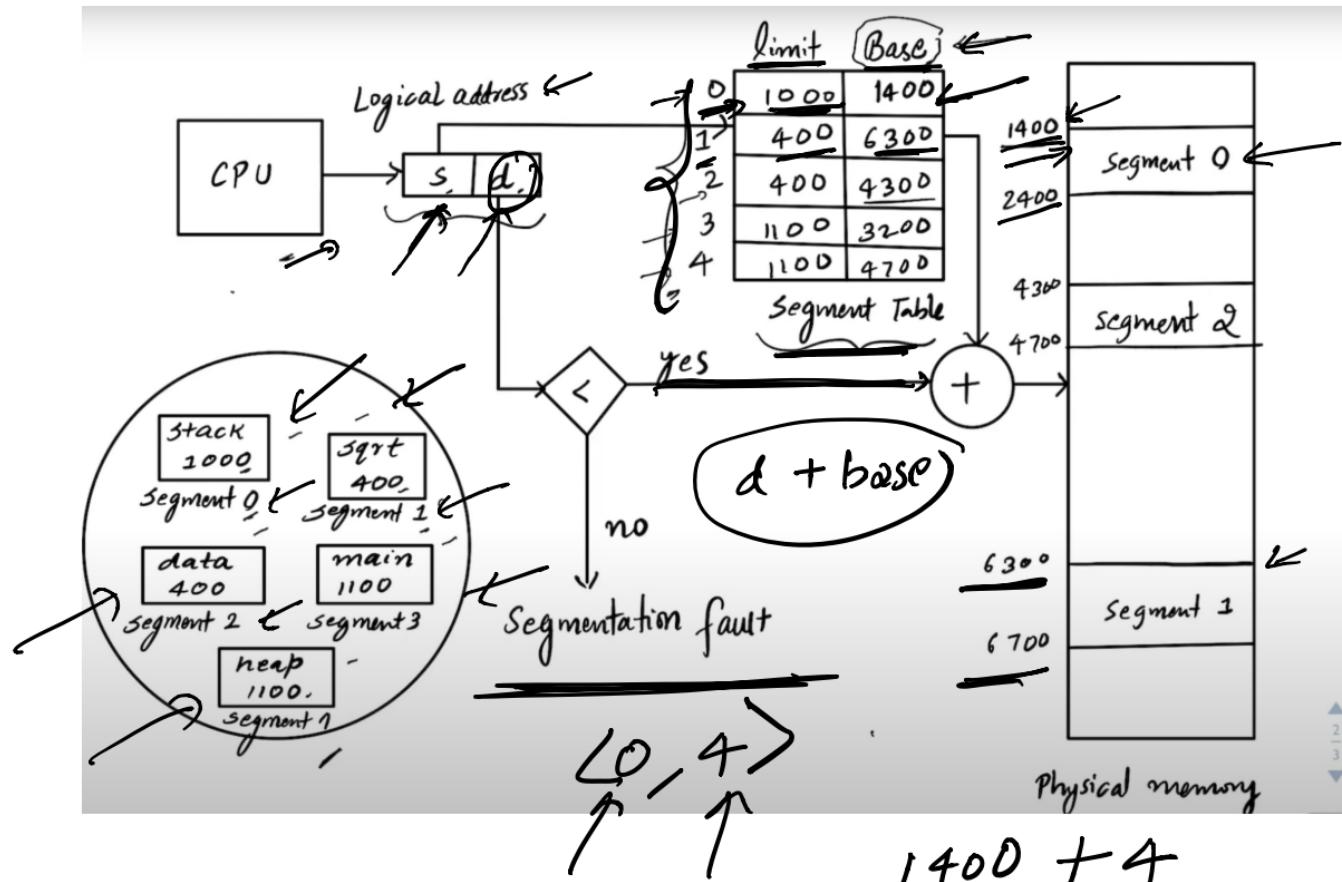
page 0	a	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
	b					f	g	h	i	r							
	c																
	d																
page 1	e	4	5	6	7												
	f																
	g																
	h																
page 2	i	8	9	10	11												
	r																
page 3		12	13	14	15												
Logical address																	



SEGMENTATION TECHNIQUE

- Users views a program collection of logical segments
 - C program ✓
 - Code segment (main(), sqrt()) ✓
 - Data segment ✓
 - Heap segment ✓
 - Stack segment ✓
 - Segment for C library ✓
- { • Segmentation is a memory management technique that supports the user view of memory.

- Logical address -> Collection of segments → *page*
- Segment name and length →
- Logical address = segment name + offset ←
pag ~



STRUCTURE OF PAGE TABLE

Hierarchical Paging



- Logical address space = 2^{32} or 2^{64}
- Page size = 4KB or 8KB
- Page table size itself can be huge and will need to be broken down
- Divide the page table into in smaller pieces



TWO-LEVEL PAGING ALGORITHM

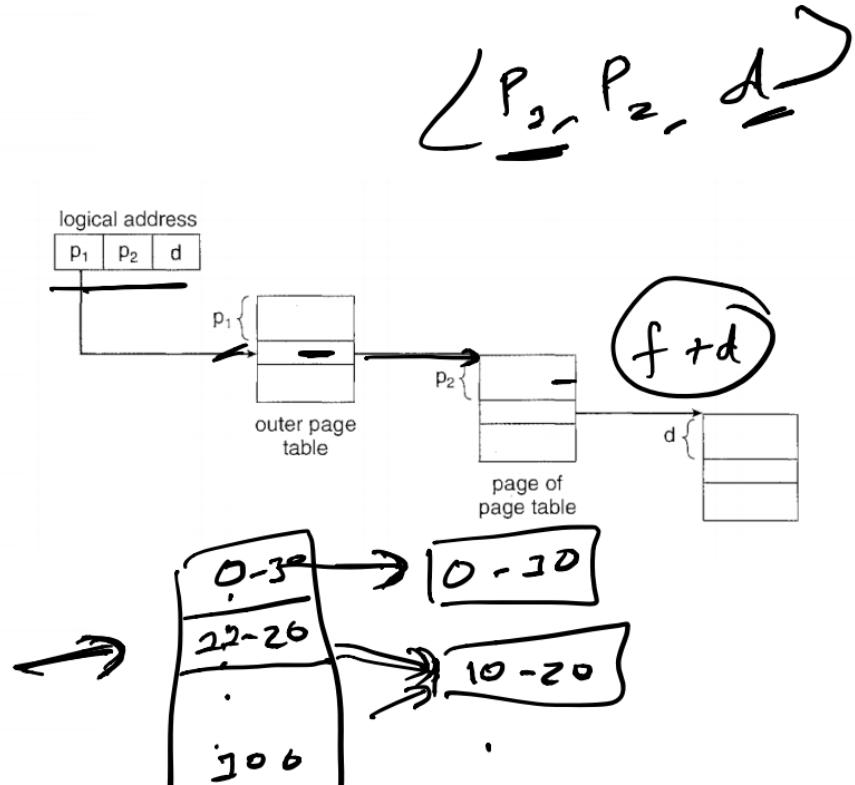
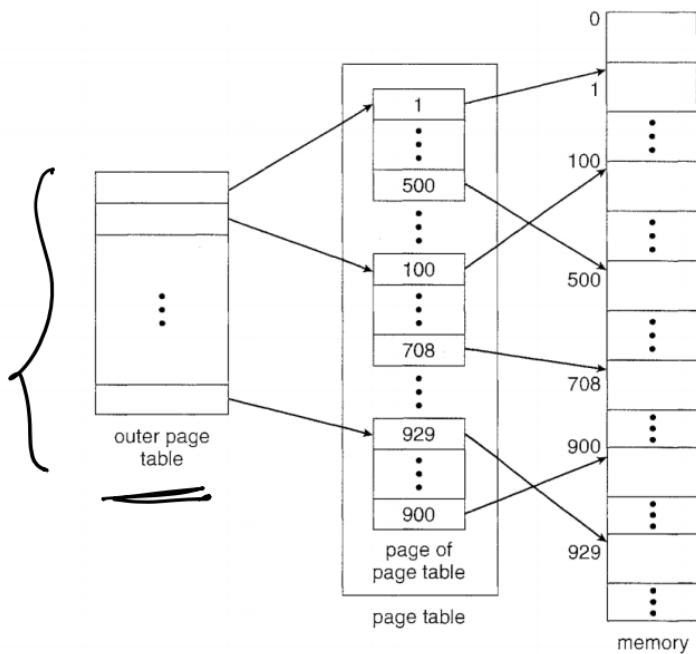
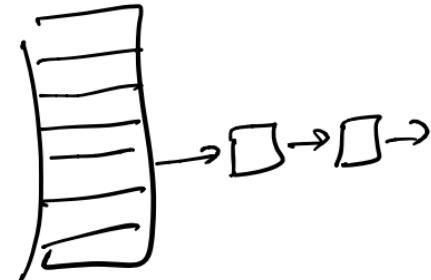


Figure 8.14 A two-level page-table scheme.

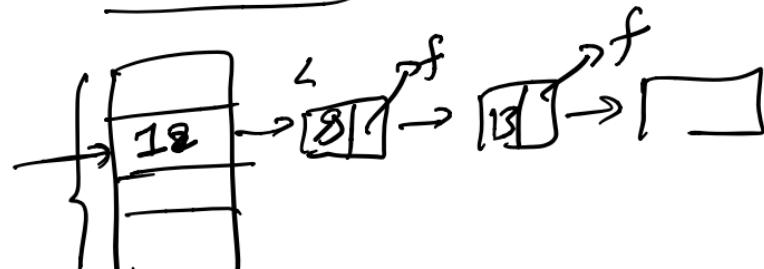
Hashed Page Tables

90 - 100

- Hashing using chaining concept
- Page number is hashed ✓
- Each node in link list has following
 - Page number ✓
 - Frame number ✓
 - Pointer to next element ✓
- Page table implementation as an array need contiguous memory

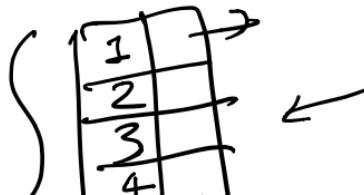


log₂ n / 1.4



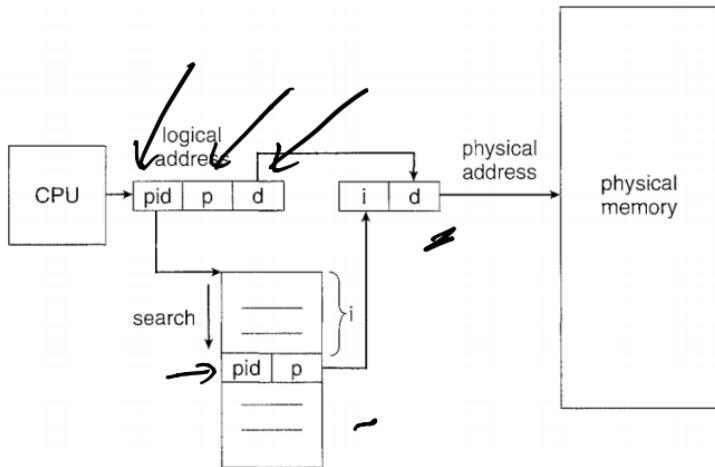
Inverted Page Tables

- Each process has its own page table ✓
- Each page table may consist of million of entries
- These tables may consume large amounts of physical memory just to keep track of how other physical memory is being used.
- In inverted page tables, we have one entry for each frame.
- Size of inverted page tables depend on the number of frames
- We need only one table now



~~ST~~

$S \rightarrow S$



<process-id, page-number, offset>.

①

Figure 8.17 Inverted page table.

$\langle pid, pag \rangle \rightarrow$

VIRTUAL MEMORY MANAGEMENT

TECHNIQUES

- Normal paging – all instructions should be in memory before program can start
- Do we really need to load the whole program at once?
 - Error handling code are used very less ✓
 - Arrays -> allocated more memory than we use 10^5
 - Certain functions are used very less ↙ 10^6
- Even if entire program is needed, it will not be needed at the same time.
- What is the maximum size of the process in simple paging?



VIRTUAL MEMORY, VIRTUAL ADDRESS SPACE

- Virtual memory management technique -> size of the process can be larger than size of main memory
- Disk is also used to store pages of the process.
- Virtual address space -> logical address space
- Virtual memory = concept of logical memory + using disk for storing pages
- When we say virtual, we mean that disk is being used for storing pages.
- Logical address is called virtual address as they are also not real.

3GB

1M

4GB

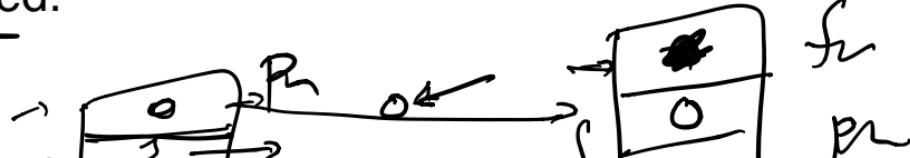


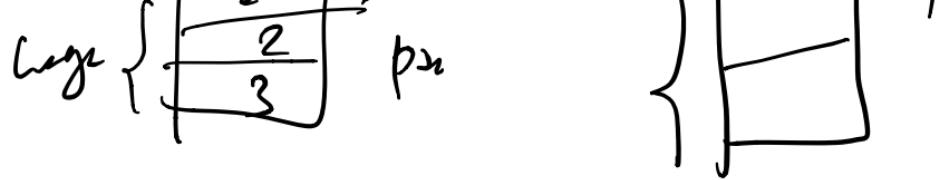
1GB



DEMAND PAGING IN OPERATING SYSTEM

- Load the page only when they are needed -> demand paging
- When a process is created, some pages are loaded into main memory by pager (frame allocation algorithms)
- CPU starts executing the program. Two things happen.
 - Page is their in main memory ✓
 - Page is not there in main memory but on disk - page fault ↙
- In page fault case, the kernel has to bring the page from disk to main memory.
- Pure demand paging -> no pages are loaded initially. When page fault happens, then a page is loaded.



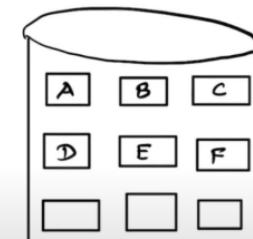


- How does kernel know about page fault? ✓
- An extra bit called valid-invalid bit is used

0	A	-
1	B	
2	C	
3	D	
4	E	
5	F	
6		

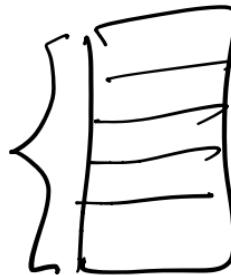
0	4	v
1	i	
2	6	v
3	i	
4	i	
5	g	v
6	i	

0
1
2
3
4
A ↙
5
6
C
7
8
9
F
10



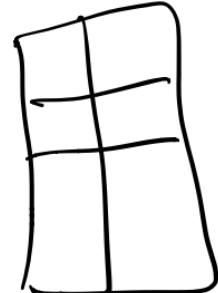
- Valid ✓

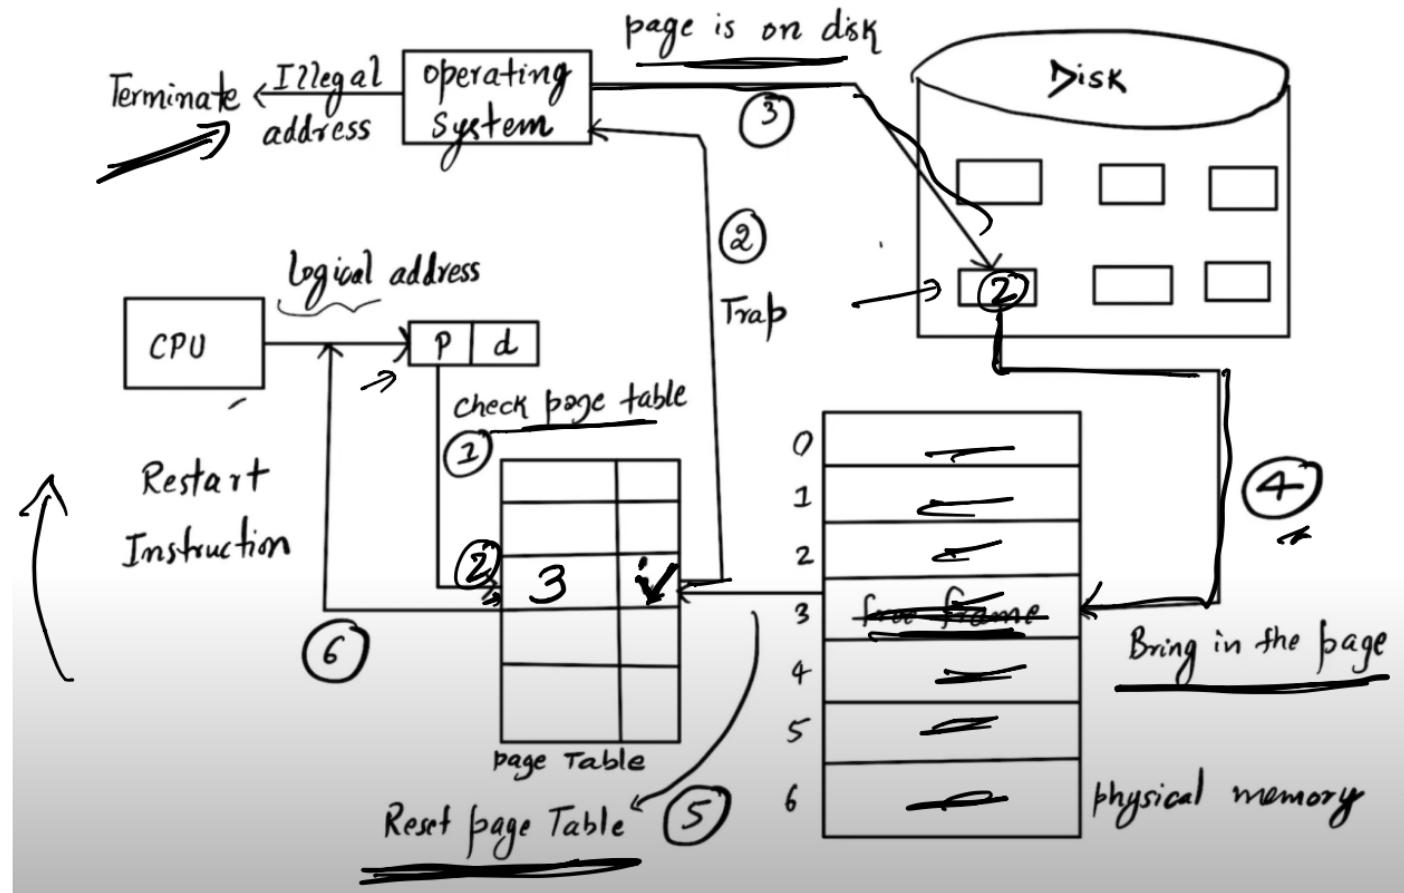
- Page is legal ✓
- Is in the memory ↗



- Invalid ↙

- Page is not legal ? Separate table in PCB keeps track of it
- Page is legal but not in main memory ?

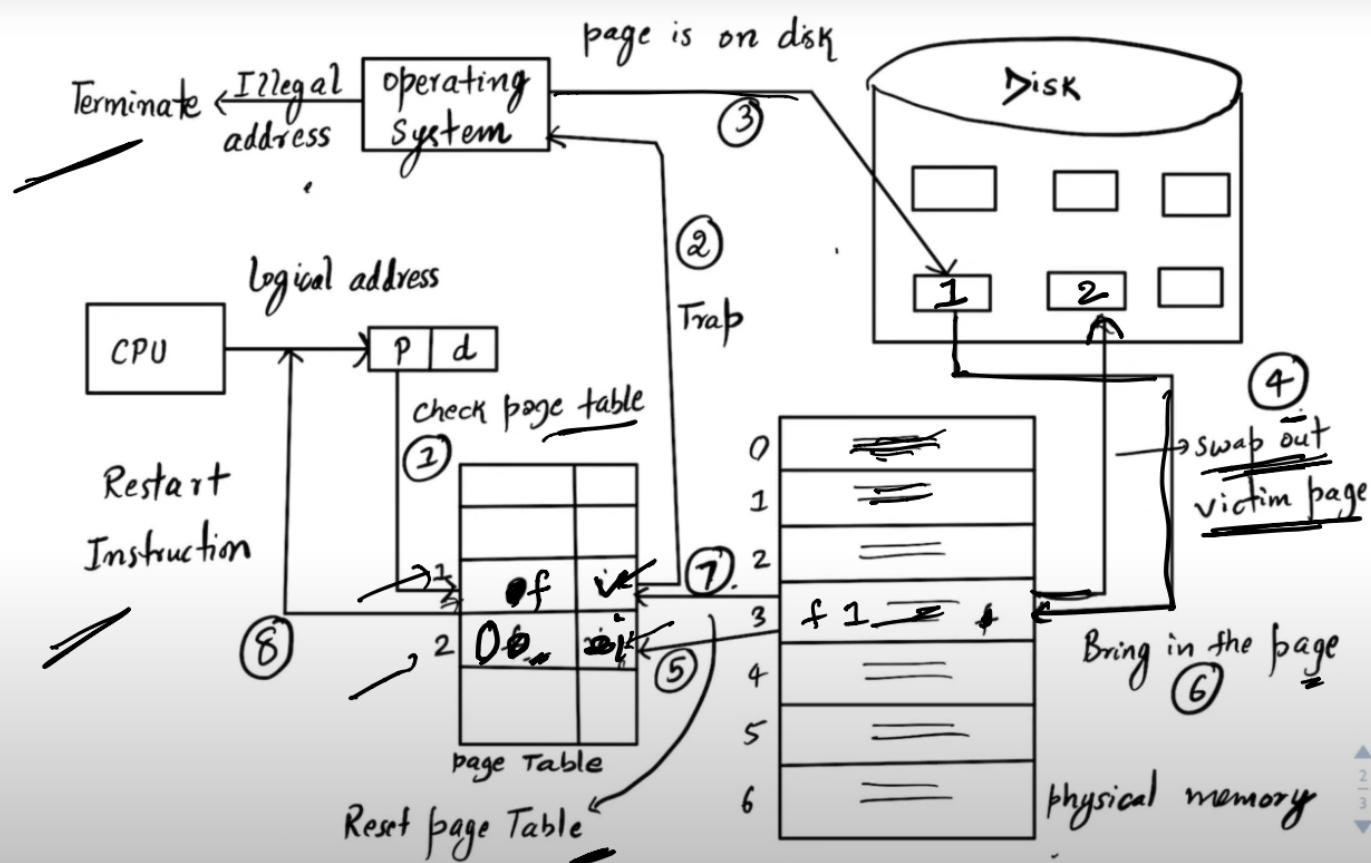




PAGE REPLACEMENT TECHNIQUES

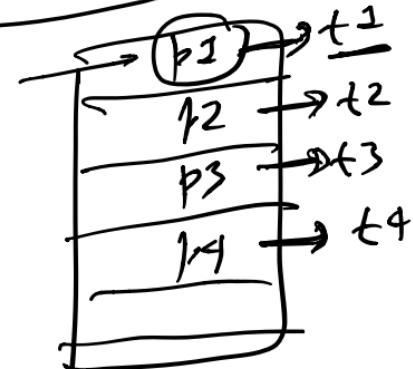
- What if all the frames are full when a page fault happens?
- Use page replacement technique to replace a page in the main memory

- - Find the location of the desired page on the disk
 - Find a free frame:
 - If there is free frame, use it
 - If there is no free frame, use a page-replacement algorithm to select a victim frame
 - Write the victim frame to the disk and change the page and frame tables accordingly
 - Read the desired page from the disk into the newly freed frame and change the frame and page tables
 - Restart the user process



FIFO PAGE REPLACEMENT ALGORITHM

- Oldest page is replaced first ✓
- First in, first out ✓
- If number of frames is increased, one would expect page faults to go down.
- Using FIFO, it's not always true. This is called Belady's anomaly.
- Example - <https://youtu.be/a2yJKKh48a4>



OPTIMAL PAGE REPLACEMENT ✓

- An algorithm with lowest page-fault rates
- Replace the page that will not be used for longest period of time
- Requires looking into future (SJF)
- Not used in OS
- Used for comparing algorithms.
- Example - <https://youtu.be/1JryPF1Ajsc>





LRU PAGE REPLACEMENT TECHNIQUE

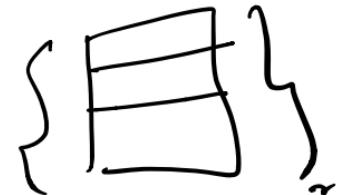
- Replace page from past that has not been used for longest period of time.
- Optimal page replacement looking backwards
- Used in most modern OS
- Doesn't suffer from Belady's anomaly
- Example - <https://youtu.be/uI3dFPkvU-I>

THRASHING

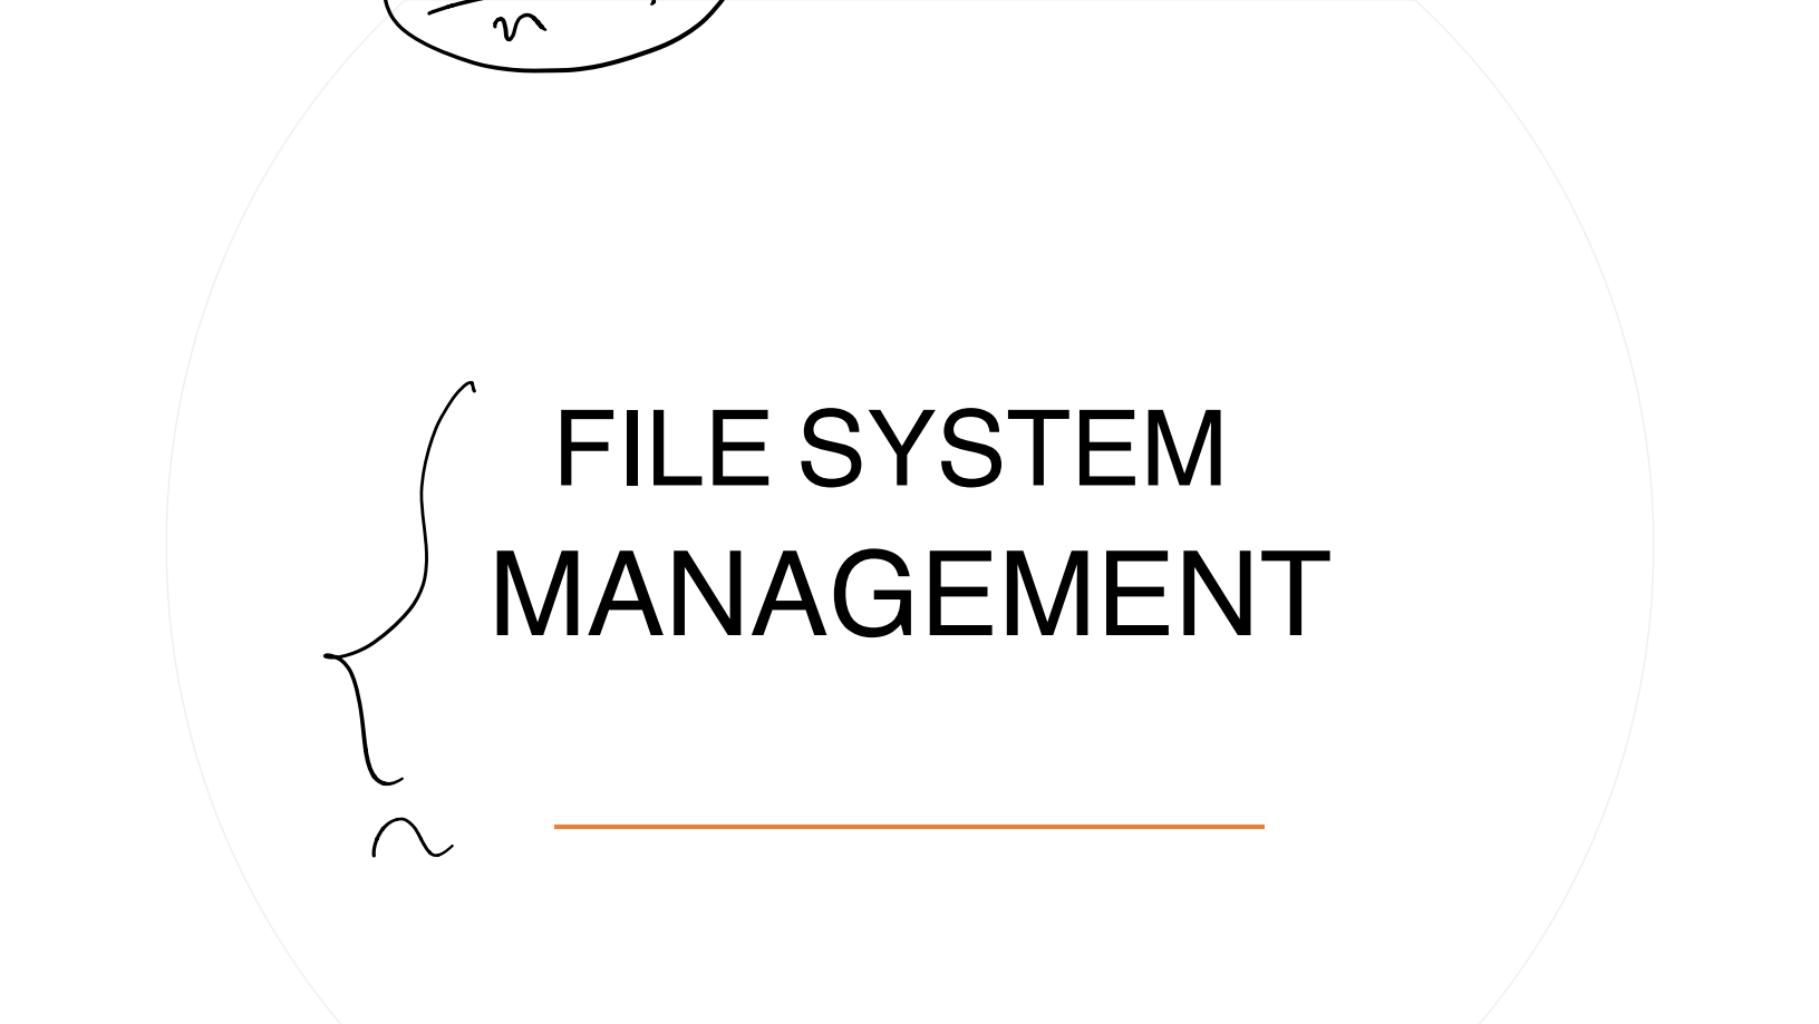
3/6

- More time is spent in swapping pages than doing actual work
- If more time is spent in swapping, CPU will be idle.
- Global page replacement -> Kernel sees that CPU are idle. It brings more process. Leads to more thrashing and a cycle.
- Local page replacement -> thrashing will not happen but local page replacement has some disadvantages.
- Equal frame allocation vs proportional frame allocation

m fr → n pag
m frame



10 KB



FILE SYSTEM MANAGEMENT

