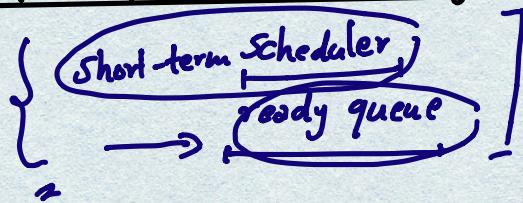


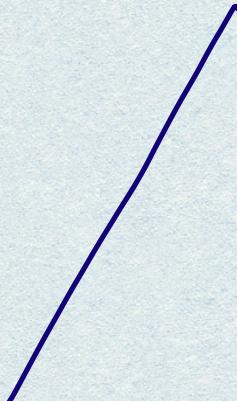
Preemptive vs Nonpreemptive Scheduling

Preemptive Scheduling



- In preemptive scheduling, each process gets a slice of CPU's time to run. 10ms
- The running process can be forcibly removed from the CPU under following conditions:
 - 1) Timeslice expires \rightarrow (running \rightarrow ready)
 - 2) An interrupt occurs on the CPU \rightarrow (running \rightarrow ready)
 - 3) A higher priority process enters the ready queue.

$\begin{cases} \rightarrow \text{NEW} \rightarrow \text{READY} \\ \rightarrow \text{WAITING} \rightarrow \text{READY} \end{cases}$



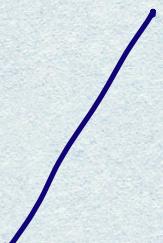
Nonpreemptive Scheduling

- In nonpreemptive scheduling, the running process cannot be forcefully removed from the CPU.
- A running process holds the CPU till :-

1) gt terminates

2) gt switches from RUNNING to (WAITING state)
on its own because of :

- a) I/O request → Wait()
- b) Waiting for the child to complete → Fork()
- c) If the mutex or the semaphore is occupied.



Examples :-

Nonpreemptive: Microsoft Windows 3.x

Preemptive: All windows version after windows 95,
Mac OS, Linux

Cost of Preemptive scheduling :-

1) sharing data between processes

- When one process is updating the data, it gets preempted.
- If some other process tries to read this data, it will find this data in inconsistent data.
- Different mechanisms like spin lock / mutex were developed to coordinate access.



2) Overhead of guarding shared data in interrupts

- whenever an interrupt happens, an interrupt handler is called.
- what if the same interrupt happens on two CPU's and the interrupt handler shares data ?
- Need to use mutex to protect critical section .

3) Increases the complexity of os design

- what happens when the kernel is changing some important data (I/o queues) and gets preempted in middle ?
- what if the kernel again tries to modify this data on behalf of some different process ?

