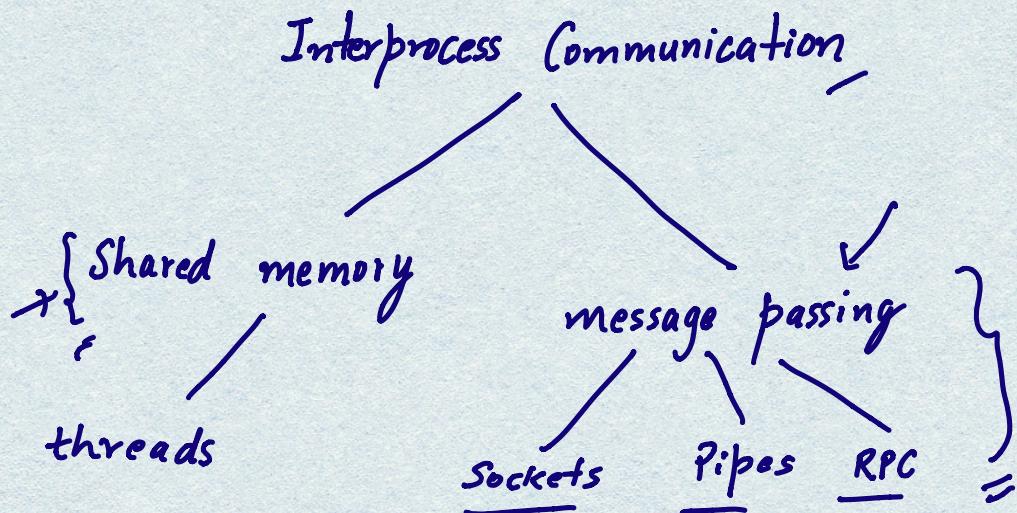


Process Synchronization



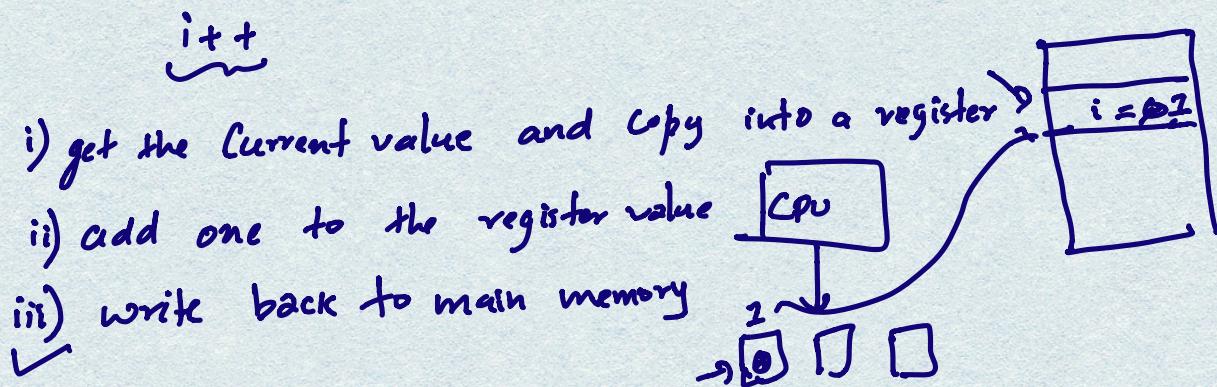
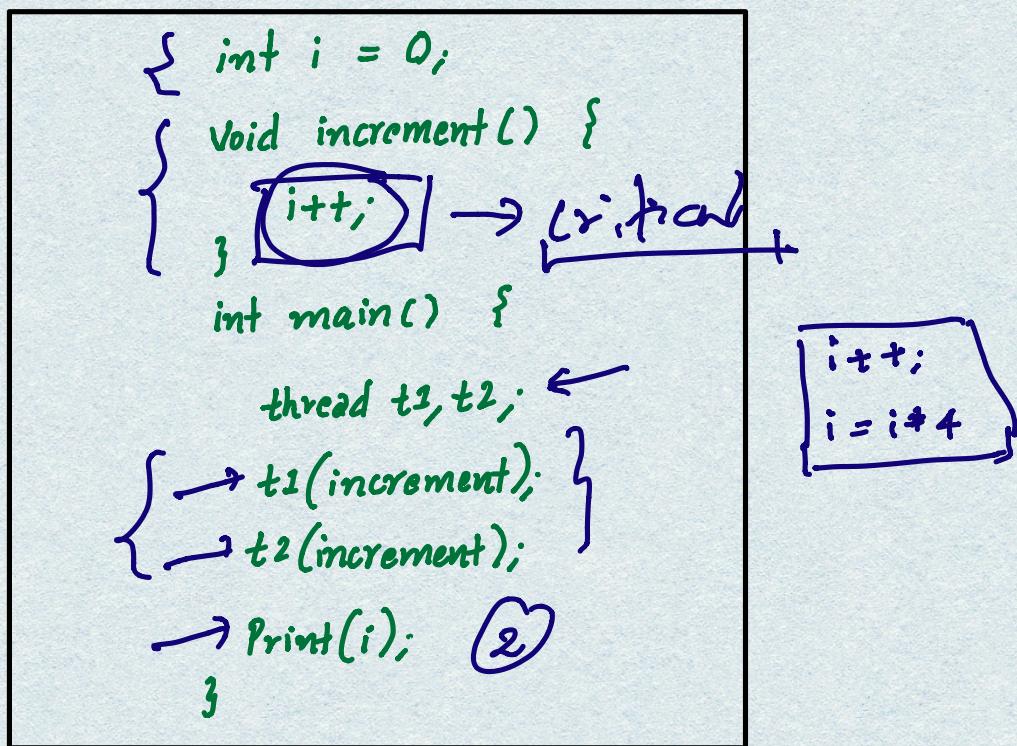
Challenges with shared memory

Concurrent access to shared data may lead to data inconsistency.

- * Discuss various issues that can arise when shared data is accessed concurrently.
- * Also, discuss solutions to these issues.

Race Condition

* We will now look at how Concurrent access to Shared data can lead to data inconsistency.



	T1	T2
t_1	get $i(0)$	
t_2	increment $i(0 \rightarrow 1)$	
t_3	write back $i(1)$	
t_4	- — - .	get $i(1)$
t_5	- — - .	increment $i(1 \rightarrow 2)$
t_6	. - - .	write back $i(2)$

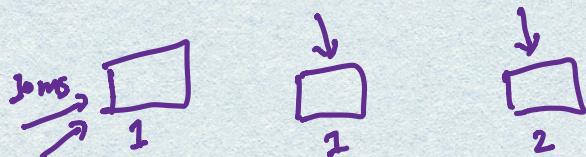
	$\rightarrow T_1$	$\rightarrow T_2$
t_1	get $i(0)$	\rightarrow get $i(0)$
t_2	increment $i(0 \rightarrow 1)$
t_3	increment $i(0 \rightarrow 1)$
t_4	write back $i(1)$	- - - .
t_5	— - - .	write back $i(1)$
		$\Rightarrow \boxed{1}$

Race Condition

When several threads tries to access and manipulate the same data concurrently and the outcome of execution depends on the order in which access takes place these type of situations are called race conditions.

* Each thread is racing for the access to the shared data.

* Note the difference between Concurrent and parallel execution.



Different ways race conditions can happen :

• Interrupts

• Kernel preemption

• A task in kernel sleeps

• Parallel processing → symmetrical multiprocessing



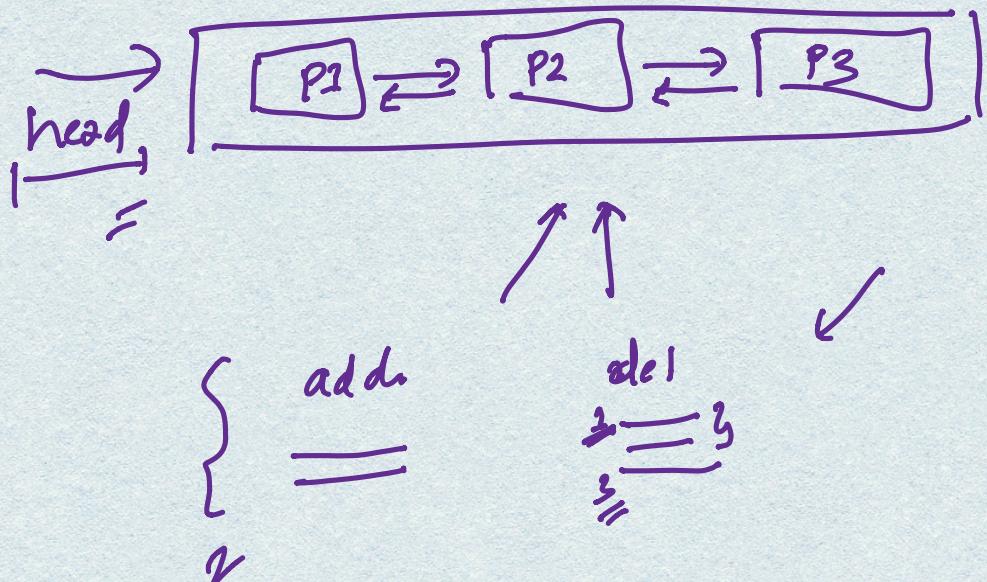
→ Concurrent execution

```

int total = get_bank_balance() → db
int withdrawal = get_withdrawal_amount()
if (total < withdrawal) {
    error("you don't have sufficient balance")
    return -1
}
total -= withdrawal
update_total_balance(total)
split_out_money(withdrawal) ←

```

Joint holder atm
2500



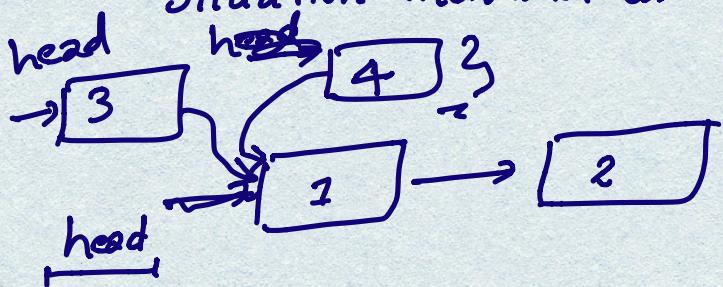
Critical Section Problem and its solutions

- Assume we have n processes or threads.
- And a Common Segment of code shared between each of these processes.

} if multiple processes tries to execute this Common section of code and it leads to data inconsistency.

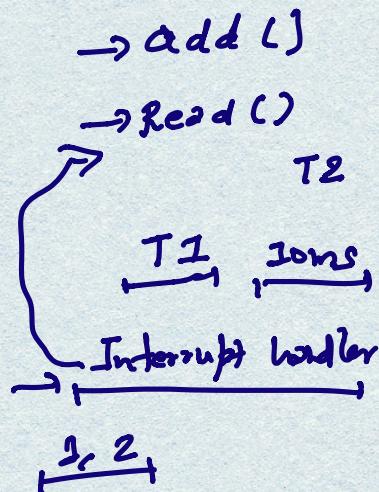
Common Section of code → Critical Section

Situation mentioned above → Race Condition



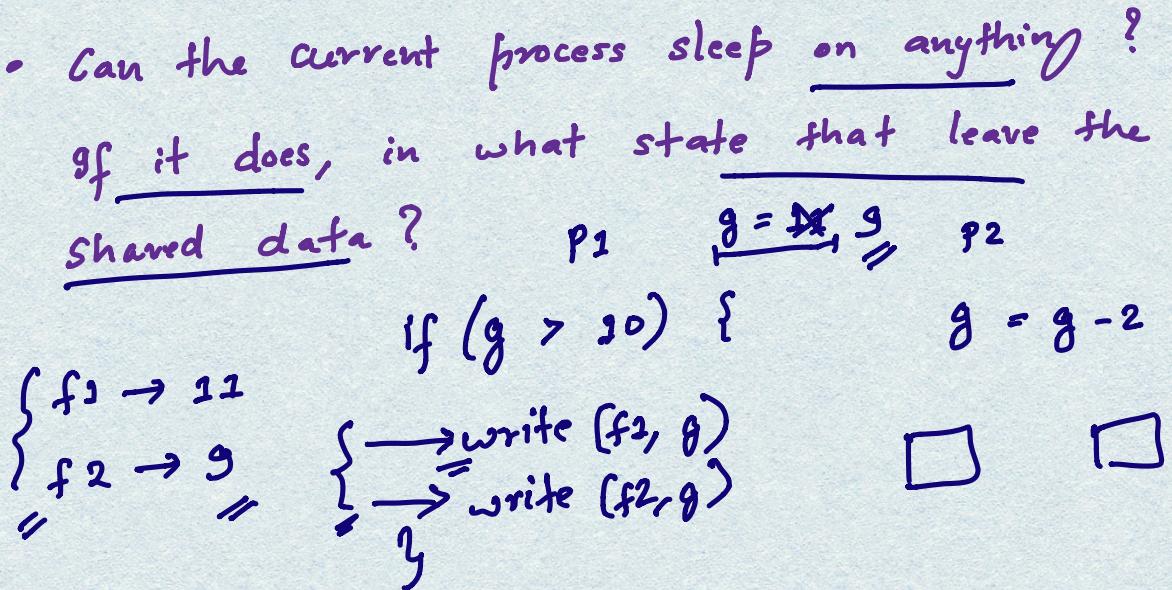
newnode = create()

newnode → next = head
head = newnode



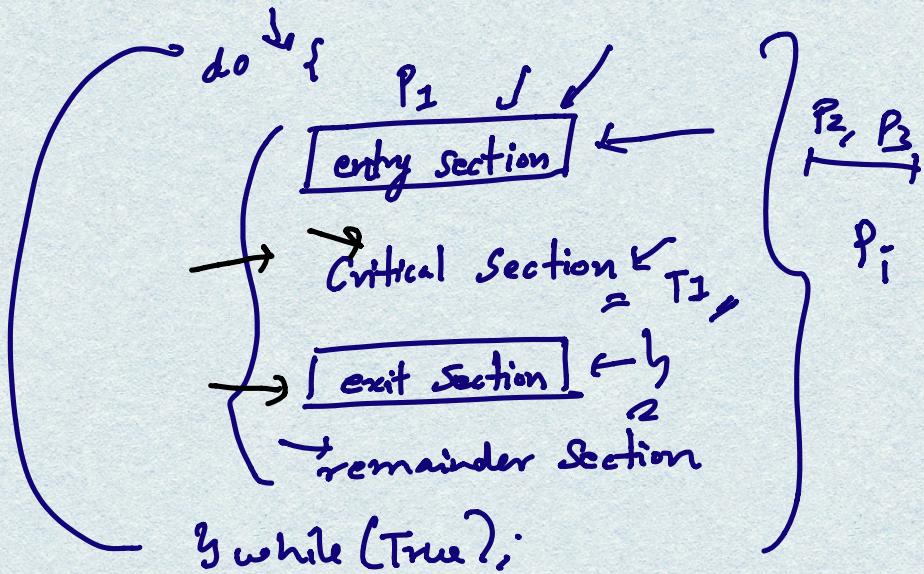
Ask yourself these questions before deciding if your code can lead to race condition or not.

- Is the data global? Can other threads access it? T₁
- Is the data shared between process context and interrupt context?
- If the process is preempted while accessing this data, can the newly scheduled process access it?
- What happens if this function is called on another processor?
- Can the current process sleep on anything? If it does, in what state that leave the shared data?



Critical Section Problem

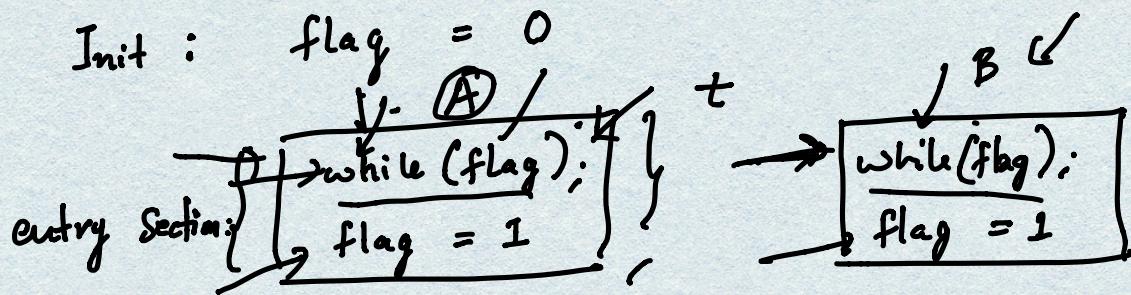
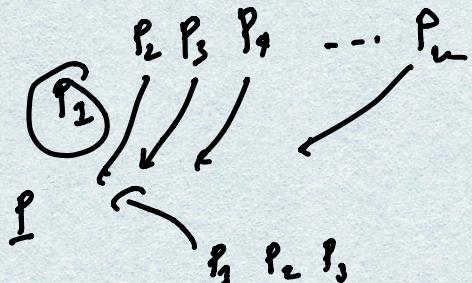
The critical section problem is to design a protocol that the processes can use to cooperate or that doesn't lead to race condition.



```
main() {  
    T3, T2,  
}
```

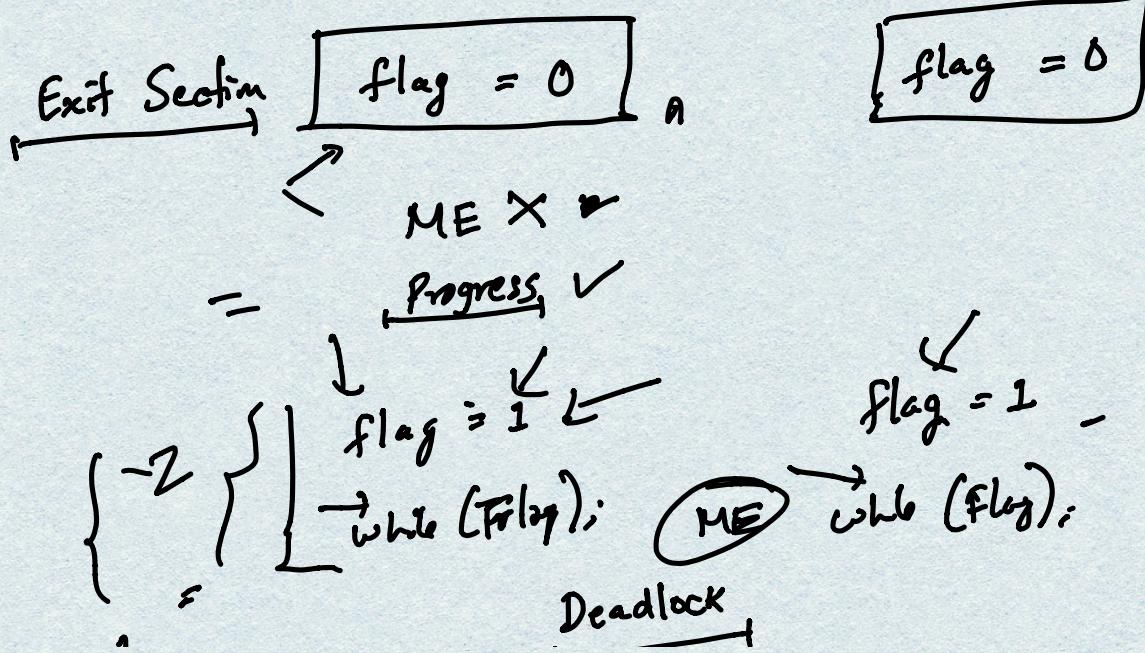
Requirements that a solution to CS problem
must satisfy : —

- i) Mutual exclusion
- ii) Progress
- iii) Bounded waiting



Critical Sect.

CS



✓ T

Init: Flag A = Flag B = False

Entry:

\rightarrow while (Flag B);

Flag A = True

while (FlagA);

Flag B = True

+ Cl

→ 5

Ex4 : $\text{Flag A} = \text{false}$

FlagB = false

→ ME X

→ Prog, es ✓

Entry: { \overline{A} $\overline{\overline{\overline{B}}}$ }
 $\Rightarrow \text{flag A} = \text{True. } \tau$
 $\Rightarrow \text{while } (\text{flag B});$ $\overline{\overline{\overline{\text{P}}}}$

$\rightarrow \text{flag } B = \text{True}$

Deadlock

$\neg y \ A = \text{false}$, ✓

Bounded waiting

Tent taking

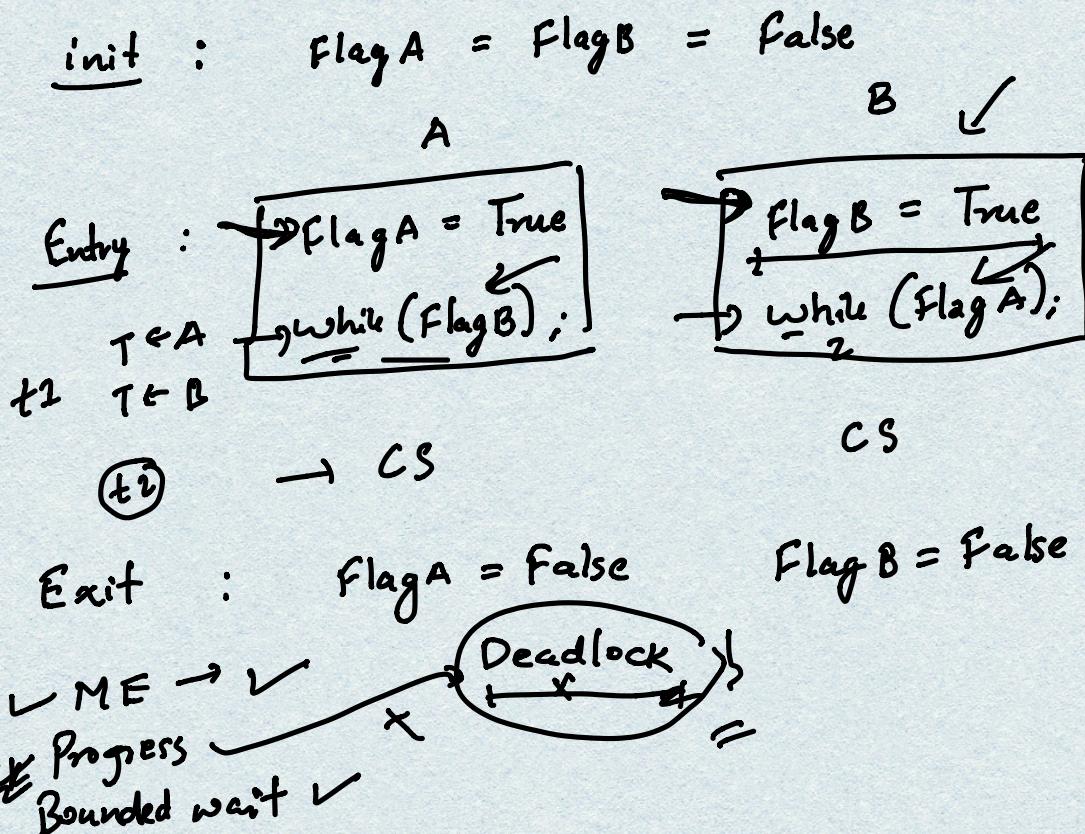
4

2

Peterson's Solution for Critical Section Problem

- ① Using two lock variables
- ② Turn-taking approach

Two lock Variables



Turn Taking approach

init:

Turn = A // B

Deadlock X

Entry :

$\left\{ \begin{array}{l} \rightarrow \text{while } (\text{turn} \neq A); \\ \text{turn} = A \end{array} \right.$
 CS

$\rightarrow \text{while } (\text{turn} \neq B);$
 $\text{turn} = B$
 CS

Exit :

$\left\{ \begin{array}{l} \text{ME } \checkmark \\ \text{Progress } \checkmark \\ \text{Bounded wait } X \end{array} \right.$
 $\xleftarrow{\text{A, B}}$

$\left\{ \begin{array}{l} \text{turn} = A \\ \text{while } (\text{turn} \neq A), . \end{array} \right\}$
 $\left\{ \begin{array}{l} \text{turn} = B \\ \text{while } (\text{turn} \neq B) \end{array} \right\}$

Init : $\text{flag A} = \text{flag B} = \text{false}$

$\therefore \text{turn} = A // B$

(X) A

Entry :

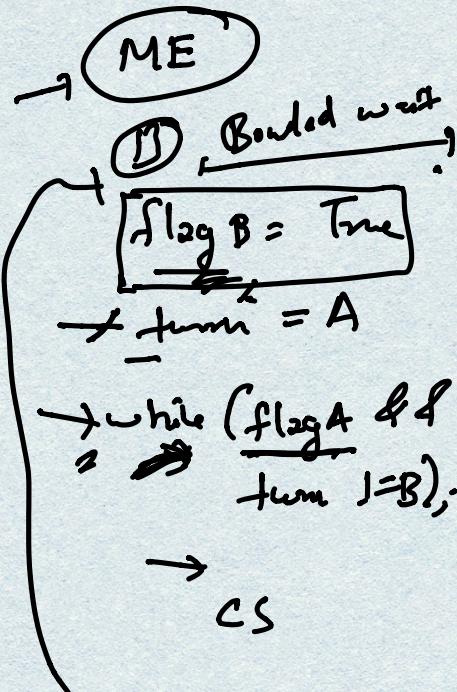
+ A, B

$\left\{ \begin{array}{l} \text{if } \text{flag A} = \text{True} \\ \rightarrow \text{turn} = B \\ \rightarrow \text{while } (\text{flag B} \& \& \\ \quad \quad \quad \text{turn} \neq B); \end{array} \right.$

$\boxed{\text{flag A} = \text{flag B} = \text{True}}$

=

CS



Ex. 1:

$\rightarrow \text{flag A} = \text{false}$

$\rightarrow \text{flag B} = \text{false}$

i, r

$\rightarrow \text{flag}[2] = \{0, 1\} \rightarrow \text{flag}[0] = \text{flag}[1] = 0$

$\rightarrow \text{turn} = A // B$

Limitations of Peterson's Solution :-

- (1) spinlock - wastes cpu cycles { N
 - (2) Doesn't work with unknown number of processes
 - (3) Doesn't work well on modern operating systems which have multiple cores. }

Init : $\text{flag}[i] = \text{flag}[j] = \text{False}$
 $\text{turn} = j // i$

Entry : $\rightarrow \text{flag}[i] = \text{True}$ $\xrightarrow{\text{store}}$ $\text{flag}[j] = \text{True}$ $\xrightarrow{\text{store}}$

A B

$\text{turn} = y$ $\xrightarrow{\text{load}}$ $\text{turn} = i$ $\xrightarrow{\text{load}}$

$\rightarrow \underline{\text{while}(\text{flag}[j] \& \text{turn} == j)}$ $\xrightarrow{\text{turn} == j};$ $\xrightarrow{\text{turn} == i};$

Critical Section Critical Section

Exit: $\text{flag}[i] = \text{False}$ $\text{flag}[j] = \text{False}$

{ Load \rightarrow It loads the value from main memory to
 CPU registers
 Store \rightarrow It stores the value from CPU registers
 to main memory.

